


• Render Post

Turn raw D5 renders into marketing-grade images. User guide.



What's in this guide

01 What Render Post does

02 Before you start

03 First launch and projects

04 The workflow, step by step

05 Settings explained

06 Versions and angles

07 Video: frames, clips, reels

08 Where your files go

09 Getting good results

10 Character

11 Troubleshooting

12 Updates and versions

01 What Render Post does

Render Post takes a folder of raw renders and gives each one a post-production pass with an AI image model. It does not re-render anything and it is not a Photoshop replacement. It is the step after the render: better light, more believable materials, real-looking people and vegetation, a considered grade.

The important part is how it gets there. For every image, an AI art director looks at the render and writes a bespoke enhancement prompt for that image. You read the prompt, change anything you like, and only then is anything sent to be enhanced. Every result is kept as a numbered version, so you can compare, go back, and pick the one you want.

Everything happens in a browser tab on your own computer. The only thing that leaves your machine is the render being processed and its prompt, which go to **fal.ai** using your own account.

Two things grow from there. **Angles** turns one finished image into new stills of the same space from other camera positions. **Video** animates any frames you choose: a short clip per frame with your image as the first shot, a reel that stitches those clips together with music, or a single continuous take.

One sentence version. Drop the exe into a folder of renders, double-click, review the prompts it writes, press Enhance, then switch to Video, choose frames, and make clips.

02 Before you start

What you need

- A Windows PC and a modern browser (Chrome, Edge or Firefox).
- A **fal.ai** account with a payment method. Create one at fal.ai, then make an API key at fal.ai/dashboard/keys. You'll paste it once on first launch.
- Renders as PNG, JPG or WebP. Any size; large ones are downscaled to 3840px on the long edge before upload.

What it costs

You pay fal directly for each image generated. Writing the prompts costs fractions of a cent per image. The enhancement itself is where the money goes:

MODEL	PER IMAGE	NOTES
GPT Image 2	Token priced, typically \$0.10 to \$0.40	Strongest at photoreal materials and people. Cost rises with quality and size.
Nano Banana Pro	\$0.15 · 4K doubles it	Deeper reasoning, tends to respect geometry better.
Nano Banana 2	\$0.08 · 2K x1.5 · 4K x2	Fastest and cheapest. Good for quick passes and variations.

Video is a different order of spend. Seedance bills per second of output:

VIDEO MODEL	PER SECOND	A 5-SECOND CLIP	NOTES
H3 Max	\$0.05 at 480p · \$0.08 at 768p	\$0.25 to \$0.40	The default. Top of the image-to-video leaderboard, renders in seconds, native audio.
Kling 3.0 Pro	\$0.14 with audio · \$0.11 silent	about \$0.70	Cinematic, around 1080p, tolerant of people in frame.
Seedance 2.5	\$0.22 at 480p · \$0.47 at 720p	\$1.10 to \$2.40	The only model for a single take across frames. Refuses frames with realistic people.

fal rates as of September 2026. Render Post shows an estimate before you press Enhance, Angles or Make, and asks you to confirm before generating video.

Installing

There is no installer. `RenderPost.exe` is the whole tool. Keep it on your Desktop or in a tools folder and copy it into any render folder when you want to use it, or double-click it from anywhere and choose the folder when asked.

Windows will warn you the first time. A blue "Windows protected your PC" screen appears because the app isn't signed with a Microsoft certificate. Click **More info**, then **Run anyway**. You only see this once per machine.

03 First launch and projects

01 Put the exe next to your renders and double-click it.


A browser tab opens within a few seconds. The first launch takes a little longer while it unpacks. If the exe is not sitting next to any images, a folder picker appears instead.

02 Paste your fal key.

The key is checked immediately. It's stored on this computer only, never inside the render folder, so you can share the folder without sharing your key.

03 Leave the tab open while you work.

Closing it stops Render Post about 30 seconds later, once nothing is in progress. There's also a Quit button top right.




The key prompt on first launch. Change it later with the Change key button.



The header: counts and current model on the left; Images / Picks / Video, the Before / Wipe / After view, the running spend, status and Quit on the right.

Projects

Each folder of renders is its own project: its versions, notes, character and video set all live inside its enhanced folder. To work on another one, press **New project** in the settings bar: recent projects are one click, or Browse to any folder. You never need to move the exe again.



New project: recent folders and Browse.

Adding renders to a project: drop files into the folder (Render Post notices on its own), or drag images onto the browser window.

04 The workflow, step by step

The default mode is **Write prompts first, let me review**. Nothing is enhanced until you say so.

01 Set style notes, if you have any.

Plain English, applied to every prompt written from now on: "cooler white balance", "dusk, no people", "autumn trees". Leave it empty and each image gets its own best guess. Notes are saved with the folder, so every project keeps its own.

02 Press Write prompts.

Each render goes to the art director model, which returns a prompt tailored to that image. This step is close to free. Cards show "prompt ready · not sent".

03 Read and edit.

The prompt sits in an editable box on each card. Change wording, delete a line, add a detail. Rewrite asks for a fresh prompt for that one image; Copy copies it.

04 Press Enhance.

The main button sends every reviewed prompt, or tick images (the box beside each name) and it becomes

Enhance N selected


. Enhance this one works on a single card. The estimate beside the button says what it will cost, and every spend asks you to confirm.

05 Compare.


Drag the orange handle to wipe between raw and enhanced. Keys **1** **2** **3** switch Before / Wipe / After for every card; **J** and **K** jump between images.

06 Keep, share, or move on.


Pick (the star) copies the version you're looking at into a picks folder. Before / after saves a side-by-side JPG. Angles makes new views of the space. Add to video sends the version to the Video section.




The settings panel at launch. In review mode the only orange button is Write prompts. More opens the less-used settings.




With two images ticked: the button, the estimate and the note follow the selection.



After Write prompts: the prompt is ready and nothing has been sent.



After an Angles run: v tabs are enhancements, a tabs are new camera positions. Angles and Add to video sit under the prompt.



While an image is enhancing. Cancel stops that one; Stop in the header cancels everything.




Change the style notes after prompts exist and Render Post flags them. Update prompts revises them without re-enhancing.

Changing notes later. Notes shape prompts when they are written, they don't rewrite existing ones on their own. When notes change, the button becomes **Update N prompts**: each prompt is revised to follow the new notes while keeping the image-specific detail it already has. Prompts you edited by hand are left alone.

05 Settings explained

SETTING	WHAT IT DOES
Style notes	Instructions to the art director. Saved per folder.
Model	GPT Image 2, Nano Banana Pro or Nano Banana 2. The controls below change with the model. A hint underneath explains strengths and price.
Quality	GPT only. Low, medium, high. Medium is the sensible default.
Output size	GPT only. Long edge from 1.5K to 4K, aspect ratio follows the render. fal caps total pixels at about 8.3 MP, so very wide images are reduced.
Resolution	Nano models only. 1K, 2K or 4K.
Variations	How many results per image in one run: 1, 2 or 4. Each becomes its own version. Multiplies the cost.
Angles per run	How many new camera positions an Angles run produces: 4, 6 or 9. Under More.
Spend warning at	The running total in the header turns amber past this figure. Under More.
Grid / Detail	Above the images. Grid shows the latest version of every image for scanning; click opens it, shift-click selects it.
show all · catalog	Beside the model dropdown. show all reveals non-recommended models; catalog opens where extra models come from (a JSON file at a URL, see the dialog for the format).
Before spending credits	Review first (default) or Enhance straight away. In the second mode new renders are prompted and enhanced as soon as they appear.



Grid view: every image at a glance, with status, version count and a star if picked.

The running spend

The header shows an estimate of what this project has cost so far, added up from fal's published rates as jobs finish. GPT Image 2 is token priced, so it counts as \$0.25 an image. It is an estimate for glancing at, not a bill; your fal dashboard is the truth.

Buttons in the settings bar

BUTTON	WHAT IT DOES
--------	--------------


Open folder	Opens the enhanced folder in Explorer.
Rescan	Looks for renders added to the folder since launch and adds them as new cards.
Change key	Replace your fal key.
New batch	Writes a fresh prompt for every image using the current notes, then waits for review. Use it for a second round of options.
New project	Switch to another folder of renders.
Export prompts	Saves every prompt, version and the hosted raw-render URL to enhanced/prompts-export.json, for use in other tools.
Enhance / Update prompts	The main action. Its label tells you what it will do next.
Stop	Appears while anything is running. Cancels everything queued or in progress.

06 Versions and angles

Render Post never overwrites a result. Every enhancement of an image becomes the next version: v1, v2, v3. On disk they are Scene 4_v01.png, Scene 4_v02.png and so on, so a batch is just everything with the same number.

- **Tabs** next to the image name switch between versions. The wipe compares the raw render against whichever one is selected, and the prompt box shows the exact prompt that made it.
- **Enhance again** sends the prompt currently in the box and adds a version. Edit the text first if you want something different.
- **Rewrite** asks the art director for a brand new prompt for this image, using the current notes.
- **Pick** (the star) marks the selected version as your choice and copies it to enhanced/picks. Press again to unpick.
- **Before / after** saves a side-by-side JPG of raw and the selected version to enhanced/compare.
- **Delete** moves the selected version to enhanced/trash and takes it out of picks and the video set. Handy for angles that missed.

The **Picks** tab in the header shows only picked versions, as a thumbnail gallery with the cards underneath. Click a thumbnail to jump to it.



The Picks tab.

Angles

Press **Angles** on a finished version and an art director plans several camera positions the image can support: a wide view from the entrance, the reverse angle looking back, a low three-quarter across the furniture, a material detail, a view through an opening. Each is generated with your current image model and lands as an a tab: a1, a2, files named Scene 4_a01.png. They appear in the Video frame grid immediately, so one hero render can become a six-frame reel.

Read angles as plausible views, not as the design. Enhancement keeps your camera and improves the pixels. Angles ask the model to imagine the rest of the room. In testing, a wide, well-lit view that shows the whole space wasted about one angle in six; a partial or steep view wasted half. Wide views hold up best, reverse angles are mostly invented, Nano Banana Pro keeps geometry best. Use Angles for reels and mood, and check anything client-facing against your model.


Angles

Make 4 new angles of lobby_view v1 with GPT Image 2 · token priced?

Cancel

Make angles

Every spend confirms in-app with the model and the cost.



The wipe. Drag the handle, or use the Before / Wipe / After switch in the header. For an angle, use After; the raw is a different view.

07 Video: frames, clips, reels


Video is its own section. Switch to it with **Images / Video** in the header, or press **Add to video** on any finished image, which adds that version to the set and takes you there. Three steps, top to bottom.



The header in the Video section. The tab shows how many frames are in the set.

Step 1 · Frames

Every image is a column with its versions stacked, enhancements and angles alike, and the raw render at the bottom. Click a version to add it to the video set, click again to remove it. The orange badge is its place in the order.



Step 1. Three frames chosen, including an angle, in the order they were clicked.

Step 2 · Make clips

The toggle at the top is the main decision: **Separate videos** (each frame becomes its own clip, the frame is shot one) or **One video through all frames** (a single continuous take touring the frames in order, always on Seedance, experimental). The line under the heading spells out exactly what you'll get with the current settings.

SETTING	WHAT IT DOES
---------	--------------

Video model	H3 Max, Kling 3.0 Pro or Seedance 2.5. The hint underneath says what each is good at and what it costs.
--------------------	---

Shots per clip	1 is one continuous camera move. 2 or 3 cut to other angles of the same space inside the clip, like a mini edit.
Resolution	Where the model offers one. Start low to iterate; make the winners again at the top setting.
Clip length	4 to 15 seconds. Cost scales with it.
Sound	Native ambient audio from the model, or silent.
Camera energy	Calm barely moves; Moderate makes a clear move across the space; Dynamic asks for bold cinematic moves with real parallax. If a model still looks static on Dynamic, that's the model.
Motion notes	Plain English guidance for every motion prompt: "slow push-ins only", "no people", "evening ambience". Saved with the folder.

STEP 2 · MAKE CLIPS

You'll get 3 separate videos, 5s each. Each one starts exactly on its frame and holds one slow camera move.

Separate videos One video through all frames

MOTION NOTES · USED WHENEVER A MOTION PROMPT IS WRITTEN

slow push-ins only, calm

VIDEO MODEL ☐ show all SHOTS PER CLIP RESOLUTION CLIP LENGTH SOUND CAMERA ENERGY

H3 Max · fal / MiniMax 1 · one continuous came 480p · iterate here 5 s Ambient audio Calm · barely moves

top-ranked image-to-video, very fast, cheapest · 480p or 768p, native audio · clips are 5s or longer

VIDEO SET

lobby_view · v1 plaza_noon · v1 lobby_view · a2

3 clips · 5s each · about \$0.75

Write 3 motion prompts


Step 2 with the video set and the outcome line.

01 Write motion prompts.

One per frame. Each names the life in the scene, says the architecture and lighting stay exactly as in the still, gives one specific slow camera move, says what the camera is not doing, and adds a line of ambient sound. That's the structure these models are documented to want. Writing again after changing notes updates the same drafts; Clear unsent removes them.

02 Read, edit, Make.


Prompts are editable. Make confirms the cost, then generates. Cancel on a clip stops that one; Stop in the header stops everything. Export prompts saves every motion prompt with its hosted frame URL for other tools.



A finished clip: play it in place, Open or Download it, edit the prompt and Make again. The Add to reel toggle on the video is how it joins Step 3.

Step 3 · Reel


Press **Add to reel** on the clips you want, in the order you want them; each toggle shows its number. Choose a crossfade. For music, press the folder button beside the Music dropdown, drop an mp3 or wav in, and it appears with a **Start at** field so you can begin at the drop rather than the intro. The reel takes the size of the largest clip. Made on your computer, free.



Step 3 with a finished reel. Open video folder opens enhanced/video.

A single take

Switch to One video through all frames and the set gains order badges with arrows. Write take prompt produces one prompt that references each frame as @Image 1, @Image 2 and describes the camera travelling between them. Takes run up to 30 seconds and use Seedance regardless of the model dropdown.



In take mode the frames are numbered; the arrows change the order the camera visits them.

People in frame. Seedance's content filter refuses reference images that contain realistic people, which describes most enhanced hospitality interiors. If a take or a Seedance clip fails with that message, use Kling or H3 Max for clips, or make a no-people batch for video. 720p (768p on H3 Max) is the ceiling for clips; Kling outputs around 1080p.

08 Where your files go

LOCATION	CONTENTS
your render folder	Your original renders. Never modified.
enhanced/	Every version: name_v01.png, name_v02.png ... and angles as name_a01.png ...
enhanced/picks/	Copies of the versions you picked. This is your deliverable set.
enhanced/compare/	Before / after JPGs.
enhanced/video/	Clips as name_v02_clip01.mp4, takes as take_clip01.mp4, reels as reel_01.mp4, plus clips.json with every motion prompt.
enhanced/prompts.json	Every prompt, per version, plus the draft for the next run. You never need to open it; it's how Render Post remembers between sessions.
enhanced/renderpost.json	This folder's style notes, motion notes and video set.
%APPDATA%\RenderPost\	Your fal key and preferences (model, quality, review mode), plus log.txt for troubleshooting. Nothing here is tied to a project.

Moving or zipping the render folder keeps everything together. Delete a version file by hand and Render Post simply stops listing it.

09 Getting good results

What the art director already knows

The prompt writer is told, for every image, to keep the architecture exactly as rendered: geometry, facade, materials, openings, camera, composition. It's told not to add, remove or swap elements, and to keep the render's own white balance unless your notes say otherwise. It then names what is actually weak in that specific image and how to fix it. Your notes sit on top of that.

Notes that work

- Mood and time of day: "late afternoon, long shadows", "overcast, soft light", "blue hour with interior glow".
- Colour: "cooler white balance", "no orange cast", "muted, editorial grade".
- Entourage: "no people", "a few people, mid-distance, no faces to camera".
- Context: "autumn trees", "wet paving", "light haze in the distance".

Keep it to two or three points. A paragraph of notes competes with what the model sees in the image.

Check before you send to a client

Image models redesign things. Even with the strongest instructions, an image model can turn brick niches into a back bar, replace a window with wall panelling, or add a person who wasn't there. It happens more with GPT Image 2 than with the Nano models, and more on interiors with a lot of fixtures. Always wipe the enhanced version against the raw before anything leaves your studio. Render Post makes this a two-second check on purpose.

If results drift too far

- Try Nano Banana Pro on the same prompt. It tends to hold geometry better.
- Add a line to the prompt naming the element that drifted: "the three brick arch niches behind the bar stay exactly as they are".
- Ask for less. "Improve lighting and materials only" is a valid prompt.

Motion that reads as architecture


- One camera move per clip. Push-in, pull-back, a slow dolly. If a motion prompt describes two moves, delete one.
- Life at natural speed. People walking, curtains, foliage, candle flicker. Anything fast reads as an effect.
- Don't ask for time of day to change inside a 5-second clip. It rarely lands and it fights the still.
- For a reel, keep clip lengths equal and crossfade at 0.6 s. Cuts work better when the camera moves are similar in speed.

A word on watermarks

Images from the Nano models carry Google's invisible SynthID watermark. It doesn't affect the picture but it is detectable. GPT Image 2 output does not carry one. Worth knowing if a client has a policy about AI-assisted imagery.

10 Character

One consistent person across a whole set instead of a different stock figure in every frame. It's under More in the image settings, off by default.



The Character block: generate from a description, or upload.

- **Generate** makes a full-length reference of the person you describe, on a plain background, which is the form image models keep consistent best. **Upload** takes your own image. The character lives with the project.
- Tick **use in enhancements and angles** and every prompt is written with both the render and the reference in view; the model places "the person shown in the second input image" once, naturally, with consistent face, hair, build and clothing. Versions made with the character carry a dot on their tab. Nano Banana Pro is the best of the three at this.
- Optional notes shape how they appear: seated, reading, never facing camera.
- Video needs nothing extra for clips: if the frame contains the character, the motion prompt keeps them consistent and gives them a small action. For a single take there's a separate tick to add them as a reference, with the usual Seedance filter caveat.

Use a generated person, or someone who agreed. A real photo of a stranger is a rights problem and tends to trip model filters. Generated characters are the safe default, and the button is right there.

11 Troubleshooting

WHAT YOU SEE

WHAT TO DO

Windows protected your PC	More info, then Run anyway. Normal for an unsigned app.
Nothing happens on double-click	Wait ten seconds; the first launch unpacks. If still nothing, check %APPDATA%\RenderPost\log.txt and send it to whoever gave you the tool. Also make sure the file is really called RenderPost.exe, not RenderPost.exe.txt.
Browser didn't open	The address is in log.txt. Paste it into any browser.
fal rejected your key	The key was revoked, mistyped or copied incompletely. Make a new one at fal.ai/dashboard/keys and use Change key.
Out of credit	Top up on fal. Render Post retries failed images with Enhance.
Rate limited	Wait a minute and press Enhance again. New fal accounts have low limits at first.
The model returned no image	The model refused the prompt. Reword it, usually by removing anything about people or brands, and try again.
Card says failed with an unfamiliar message	Enhance again once. If it repeats, send the message and log.txt to whoever gave you the tool.
A render I added isn't listed	Press Rescan. Only PNG, JPG and WebP are picked up.
Connection refused in the tab	Render Post was closed. It quits about a minute after the tab is closed, or after 30 minutes with no tab open at all. Double-click the exe again; the folder picks up where it was.
Nothing in the Video set	Add frames in Step 1, or press Add to video on an image.
H3 Max says duration or resolution is wrong	Fixed in 1.6.1. If you still see it, update.
The spend figure looks off	It's an estimate from list prices; GPT Image 2 is counted at \$0.25 an image. Check the fal dashboard for the real number.
The model returned no video	The model refused the prompt. Remove anything about people's faces, brands or text and Make again.
Seedance refused the input images	Its filter flagged realistic people. Switch the video model to Kling or H3 Max, or enhance a no-people batch for video.
Stitch is greyed out	The status line beside it says why: it needs two finished clips with Add to reel on. If it says ffmpeg is missing, the exe was built without it; tell whoever built it.
A clip plays without sound	Sound was set to Silent when it was made, or the browser muted autoplay. Clips are muted by default in the page; use the player's

sound control. The file itself has audio if Sound was on.

Old style notes appear in a new folder

They don't; notes are per folder. If you see notes, that folder was used before.

12 Updates and versions

The version number sits next to the name in the header, for example v1.8.0. Mention it when reporting a problem.

Updating is replacing the file: download the new RenderPost.exe and delete the old one. Your key, preferences and every project folder are untouched. If the person distributing the tool has enabled it, an orange "update available" link appears in the header when a newer version is out, and new models can arrive without an update at all, through the model catalog.

Keyboard reference

KEY	ACTION
1 · 2 · 3	Before, Wipe, After for every card
J · K or arrow keys	Next and previous image
drag	Move the wipe

Render Post is built for ArchViz Academy students. It runs entirely on your machine and bills only your own fal account. Prompts and renders are sent to fal.ai and its model providers for processing; check their terms for how they handle uploaded content.